

Importing Libraries

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns

from warnings import filterwarnings
filterwarnings(action='ignore')
```

Loading Dataset

```
wine = pd.read_csv("winequality-red.csv")
wine.sample(25)
```

	fixed acidity	volatile acidity	citric acid	residual sugar
591	6.6	0.390	0.49	1.70
1171	7.1	0.590	0.00	2.20
289	11.6	0.420	0.53	3.30
1424	8.3	0.260	0.37	1.40
948	8.9	0.120	0.45	1.80
623	7.9	0.510	0.25	2.90
818	7.1	0.715	0.00	2.35
558	10.9	0.530	0.49	4.60
351	9.1	0.795	0.00	2.60
683	8.1	0.780	0.23	2.60
1453	7.6	0.490	0.33	1.90
1377	5.2	0.490	0.26	2.30
627	8.8	0.600	0.29	2.20
1227	9.0	0.580	0.25	2.00

163 0.070	7.4	0.600	0.26	7.30
1535 0.075	7.0	0.550	0.13	2.20
0 0.076	7.4	0.700	0.00	1.90
1461 0.060	6.2	0.785	0.00	2.10
1083 0.072	8.7	0.420	0.45	2.40
162 0.076	7.8	0.530	0.04	1.70
1113 0.074	8.9	0.240	0.39	1.60
11 0.071	7.5	0.500	0.36	6.10
1156 0.071	8.5	0.180	0.51	1.75
504 0.077	10.5	0.240	0.42	1.80
568 0.250	9.8	0.500	0.49	2.60

	free sulfur dioxide	total sulfur dioxide	density	pH
free sulphates \				
591 0.50	23.0	149.0	0.99220	3.12
1171 0.68	26.0	44.0	0.99522	3.42
289 0.95	33.0	98.0	1.00100	3.20
1424 0.70	8.0	23.0	0.99740	3.26
948 0.76	10.0	21.0	0.99552	3.41
623 0.96	21.0	45.0	0.99740	3.49
818 0.45	21.0	47.0	0.99632	3.29
558 0.56	10.0	17.0	1.00020	3.07
351 0.83	11.0	26.0	0.99940	3.35
683 0.56	5.0	15.0	0.99700	3.37
1453 0.58	27.0	85.0	0.99706	3.41
1377 0.62	23.0	74.0	0.99530	3.71

627	5.0	15.0	0.99880	3.36
0.49				
1227	8.0	21.0	0.99769	3.27
0.72				
163	36.0	121.0	0.99820	3.37
0.49				
1535	15.0	35.0	0.99590	3.36
0.59				
0	11.0	34.0	0.99780	3.51
0.56				
1461	6.0	13.0	0.99664	3.59
0.61				
1083	32.0	59.0	0.99617	3.33
0.77				
162	17.0	31.0	0.99640	3.33
0.56				
1113	3.0	10.0	0.99698	3.12
0.59				
11	17.0	102.0	0.99780	3.35
0.80				
1156	45.0	88.0	0.99524	3.33
0.76				
504	6.0	22.0	0.99760	3.21
1.05				
568	5.0	20.0	0.99900	3.31
0.79				

	alcohol	quality
591	11.5	6
1171	10.8	6
289	9.2	5
1424	9.6	6
948	11.9	7
623	12.1	6
818	9.4	5
558	11.7	6
351	9.4	6
683	11.3	5
1453	9.0	5
1377	12.2	6
627	9.1	5
1227	9.6	5
163	9.4	5
1535	9.7	6
0	9.4	5
1461	10.0	4
1083	12.0	6
162	10.0	6
1113	9.5	6

11	10.5	5
1156	11.8	7
504	10.8	7
568	10.7	6

```
wine.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1599 entries, 0 to 1598
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	fixed acidity	1599 non-null	float64
1	volatile acidity	1599 non-null	float64
2	citric acid	1599 non-null	float64
3	residual sugar	1599 non-null	float64
4	chlorides	1599 non-null	float64
5	free sulfur dioxide	1599 non-null	float64
6	total sulfur dioxide	1599 non-null	float64
7	density	1599 non-null	float64
8	pH	1599 non-null	float64
9	sulphates	1599 non-null	float64
10	alcohol	1599 non-null	float64
11	quality	1599 non-null	int64

```
dtypes: float64(11), int64(1)
```

```
memory usage: 150.0 KB
```

Description

```
wine.describe()
```

	fixed acidity	volatile acidity	citric acid	residual sugar \
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806
std	1.741096	0.179060	0.194801	1.409928
min	4.600000	0.120000	0.000000	0.900000
25%	7.100000	0.390000	0.090000	1.900000
50%	7.900000	0.520000	0.260000	2.200000
75%	9.200000	0.640000	0.420000	2.600000
max	15.900000	1.580000	1.000000	15.500000

	chlorides	free sulfur dioxide	total sulfur dioxide
density \			
count	1599.000000	1599.000000	1599.000000
mean	0.087467	15.874922	46.467792
std	0.047065	10.460157	32.895324

```

0.001887
min      0.012000      1.000000      6.000000
0.990070
25%      0.070000      7.000000      22.000000
0.995600
50%      0.079000      14.000000     38.000000
0.996750
75%      0.090000      21.000000     62.000000
0.997835
max      0.611000      72.000000    289.000000
1.003690

```

	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	3.311113	0.658149	10.422983	5.636023
std	0.154386	0.169507	1.065668	0.807569
min	2.740000	0.330000	8.400000	3.000000
25%	3.210000	0.550000	9.500000	5.000000
50%	3.310000	0.620000	10.200000	6.000000
75%	3.400000	0.730000	11.100000	6.000000
max	4.010000	2.000000	14.900000	8.000000

Finding Null Values

```
wine.isnull().sum()
```

```

fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density           0
pH                0
sulphates         0
alcohol           0
quality           0
dtype: int64

```

```
wine.groupby('quality').mean()
```


	fixed acidity	volatile acidity	citric acid	residual sugar
quality				
3	8.360000	0.884500	0.171000	2.635000

4	7.779245	0.693962	0.174151	2.694340
5	8.167254	0.577041	0.243686	2.528855
6	8.347179	0.497484	0.273824	2.477194
7	8.872362	0.403920	0.375176	2.720603
8	8.566667	0.423333	0.391111	2.577778
density \				
quality				
3	0.122500	11.000000	24.900000	
0.997464				
4	0.090679	12.264151	36.245283	
0.996542				
5	0.092736	16.983847	56.513950	
0.997104				
6	0.084956	15.711599	40.869906	
0.996615				
7	0.076588	14.045226	35.020101	
0.996104				
8	0.068444	13.277778	33.444444	
0.995212				
pH sulphates alcohol				
quality				
3	3.398000	0.570000	9.955000	
4	3.381509	0.596415	10.265094	
5	3.304949	0.620969	9.899706	
6	3.318072	0.675329	10.629519	
7	3.290754	0.741256	11.465913	
8	3.267222	0.767778	12.094444	

Data Analysis

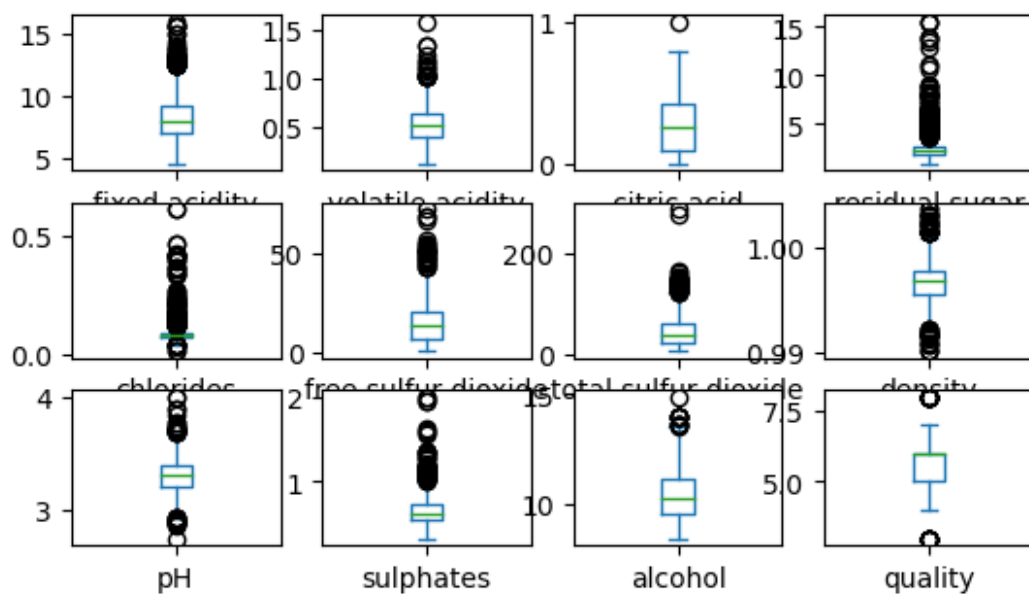
Countplot:

```
sns.countplot(wine['quality'])
plt.show()
```




```
wine.plot(kind = 'box', subplots = True, layout = (4,4), sharex = False)
```

fixed acidity	Axes(0.125,0.712609;0.168478x0.167391)
volatile acidity	Axes(0.327174,0.712609;0.168478x0.167391)
citric acid	Axes(0.529348,0.712609;0.168478x0.167391)
residual sugar	Axes(0.731522,0.712609;0.168478x0.167391)
chlorides	Axes(0.125,0.511739;0.168478x0.167391)
free sulfur dioxide	Axes(0.327174,0.511739;0.168478x0.167391)
total sulfur dioxide	Axes(0.529348,0.511739;0.168478x0.167391)
density	Axes(0.731522,0.511739;0.168478x0.167391)
pH	Axes(0.125,0.31087;0.168478x0.167391)
sulphates	Axes(0.327174,0.31087;0.168478x0.167391)
alcohol	Axes(0.529348,0.31087;0.168478x0.167391)
quality	Axes(0.731522,0.31087;0.168478x0.167391)
dtype: object	



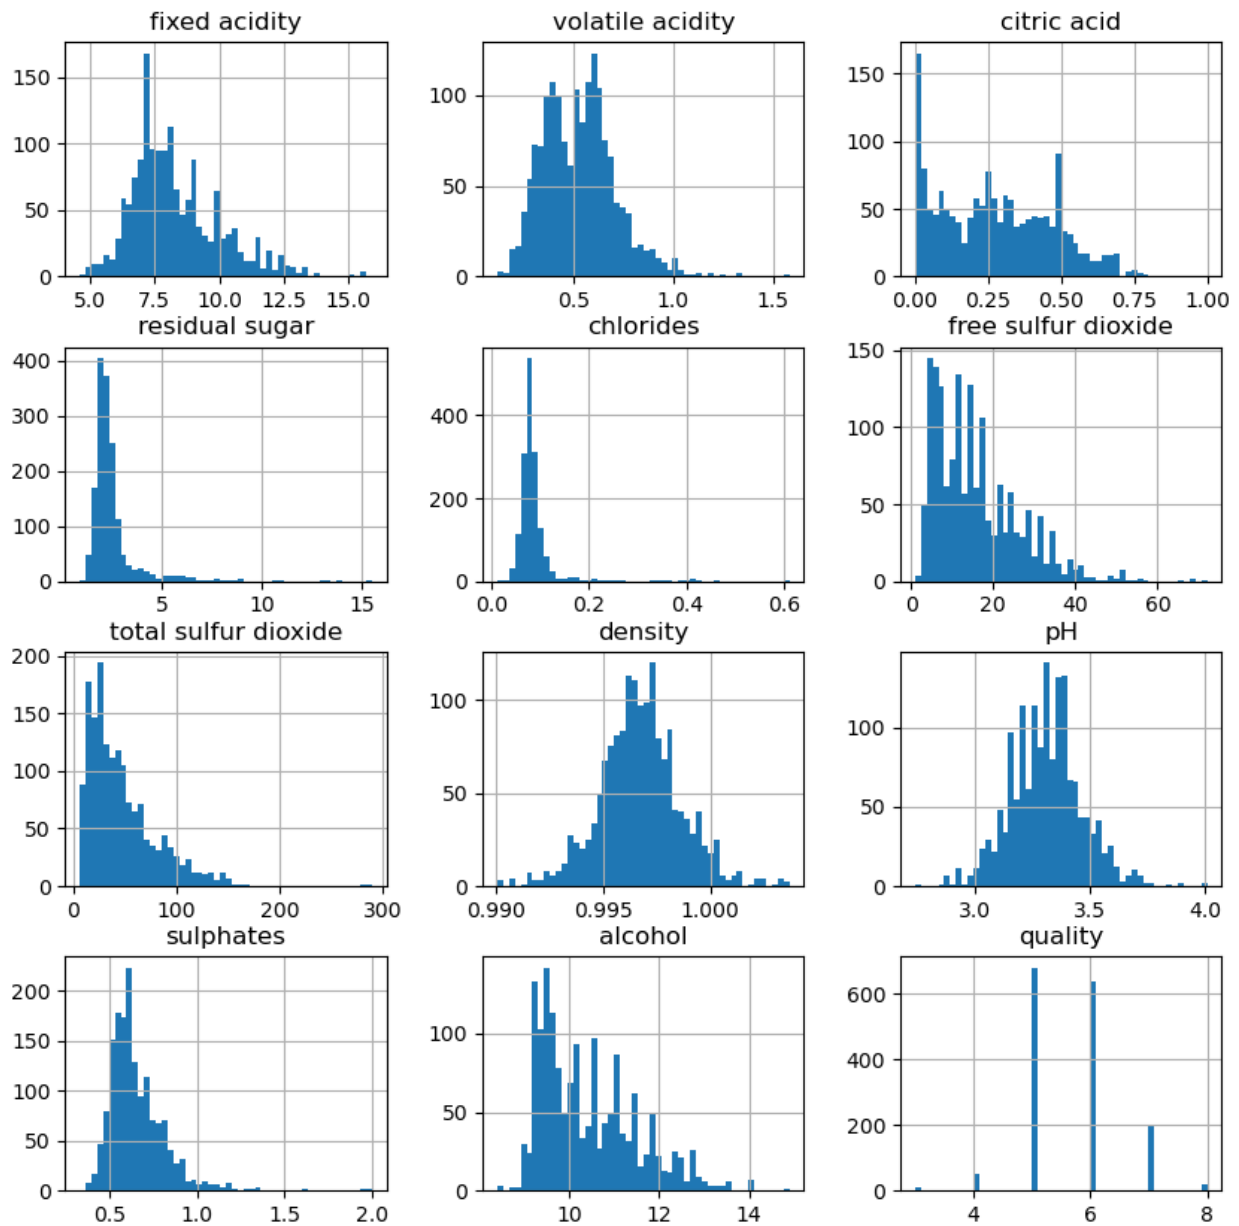
```
wine['fixed acidity'].plot(kind='box')
```

<Axes: >



Histogram

```
wine.hist(figsize=(10,10),bins=50)  
plt.show()
```



Feature Selection

```
wine.sample(5)  
  
fixed acidity  volatile acidity  citric acid  residual sugar  
chlorides \
```

1081	7.9	0.30	0.68	8.3
0.050				
471	9.6	0.54	0.42	2.4
0.081				
908	7.4	0.52	0.13	2.4
0.078				
1008	8.9	0.35	0.40	3.6
0.110				
580	12.3	0.50	0.49	2.2
0.089				

	free sulfur dioxide	total sulfur dioxide	density	pH
sulphates \				
1081	37.5	289.0	0.99316	3.01
0.51				
471	25.0	52.0	0.99700	3.20
0.71				
908	34.0	61.0	0.99528	3.43
0.59				
1008	12.0	24.0	0.99549	3.23
0.70				
580	5.0	14.0	1.00020	3.19
0.44				

	alcohol	quality
1081	12.3	7
471	11.4	6
908	10.8	6
1008	12.0	7
580	9.6	5

```
wine['quality'].unique()
```

```
array([5, 6, 7, 4, 8, 3], dtype=int64)
```

```
# If wine quality is 7 or above then will consider as good quality
wine
```

```
wine['goodquality'] = [1 if x >= 7 else 0 for x in wine['quality']]
wine.sample(5)
```

	fixed acidity	volatile acidity	citric acid	residual sugar
chlorides \				
451	8.4	0.37	0.53	1.8
0.413				
235	7.2	0.63	0.00	1.9
0.097				
1128	10.0	0.43	0.33	2.7
0.095				
304	8.4	0.65	0.60	2.1
0.112				

1276	8.5	0.40	0.40	6.3
0.050				

	free sulfur dioxide	total sulfur dioxide	density	pH
sulphates \				
451	9.0	26.0	0.99790	3.06
1.06				
235	14.0	38.0	0.99675	3.37
0.58				
1128	28.0	89.0	0.99840	3.22
0.68				
304	12.0	90.0	0.99730	3.20
0.52				
1276	3.0	10.0	0.99566	3.28
0.56				

	alcohol	quality	goodquality
451	9.1	6	0
235	9.0	6	0
1128	10.0	5	0
304	9.2	5	0
1276	12.0	4	0

See total number of good vs bad wines samples

```
wine['goodquality'].value_counts()
```

```
goodquality
```

```
0    1382
```

```
1     217
```

```
Name: count, dtype: int64
```

Separate dependent and independent variables

```
X = wine.drop(['quality', 'goodquality'], axis = 1)
```

```
Y = wine['goodquality']
```

```
X
```

	fixed acidity	volatile acidity	citric acid	residual sugar
chlorides \				
0	7.4	0.700	0.00	1.9
0.076				
1	7.8	0.880	0.00	2.6
0.098				
2	7.8	0.760	0.04	2.3
0.092				
3	11.2	0.280	0.56	1.9
0.075				
4	7.4	0.700	0.00	1.9
0.076				
...
...				

1594	6.2	0.600	0.08	2.0
0.090				
1595	5.9	0.550	0.10	2.2
0.062				
1596	6.3	0.510	0.13	2.3
0.076				
1597	5.9	0.645	0.12	2.0
0.075				
1598	6.0	0.310	0.47	3.6
0.067				

	free sulfur dioxide	total sulfur dioxide	density	pH
sulphates \				
0	11.0	34.0	0.99780	3.51
0.56				
1	25.0	67.0	0.99680	3.20
0.68				
2	15.0	54.0	0.99700	3.26
0.65				
3	17.0	60.0	0.99800	3.16
0.58				
4	11.0	34.0	0.99780	3.51
0.56				
...
...				
1594	32.0	44.0	0.99490	3.45
0.58				
1595	39.0	51.0	0.99512	3.52
0.76				
1596	29.0	40.0	0.99574	3.42
0.75				
1597	32.0	44.0	0.99547	3.57
0.71				
1598	18.0	42.0	0.99549	3.39
0.66				

	alcohol
0	9.4
1	9.8
2	9.8
3	9.8
4	9.4
...	...
1594	10.5
1595	11.2
1596	11.0
1597	10.2
1598	11.0

[1599 rows x 11 columns]

```
print(Y)
0      0
1      0
2      0
3      0
4      0
..
1594   0
1595   0
1596   0
1597   0
1598   0
Name: goodquality, Length: 1599, dtype: int64
```

Feature Importance

```
from sklearn.ensemble import ExtraTreesClassifier
classifiern = ExtraTreesClassifier()
classifiern.fit(X,Y)
score = classifiern.feature_importances_
print(score)

[0.07484542 0.10041076 0.098194    0.07704065 0.06799978 0.06684995
 0.08334131 0.08620213 0.06586187 0.10751625 0.17173788]
```

Splitting Dataset

```
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test =
train_test_split(X,Y,test_size=0.3,random_state=7)
```

Result

```
model_res=pd.DataFrame(columns=['Model', 'Score'])
```

Logistic Regression

```
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train,Y_train)
y_pred = model.predict(X_test)
```

```

from sklearn.metrics import accuracy_score, confusion_matrix
# accuracy_score(Y_test, Y_pred)
model_res.loc[len(model_res)] = ['LogisticRegression',
accuracy_score(Y_test, y_pred)]
model_res

```

	Model	Score
0	LogisticRegression	0.872917

SVC

```

from sklearn.svm import SVC
model = SVC()
model.fit(X_train, Y_train)
y_pred = model.predict(X_test)

from sklearn.metrics import accuracy_score
print("Accuracy Score:", accuracy_score(Y_test, y_pred))
model_res.loc[len(model_res)] = ['SVC', accuracy_score(Y_test, y_pred)]
model_res

```

Accuracy Score: 0.86875

	Model	Score
0	LogisticRegression	0.872917
1	SVC	0.868750

Decision Tree Classifier

```

from sklearn.tree import DecisionTreeClassifier
model = DecisionTreeClassifier(criterion='entropy', random_state=7)
model.fit(X_train, Y_train)
y_pred = model.predict(X_test)

from sklearn.metrics import accuracy_score
print("Accuracy Score:", accuracy_score(Y_test, y_pred))
model_res.loc[len(model_res)] = ['DecisionTreeClassifier',
accuracy_score(Y_test, y_pred)]
model_res

```

Accuracy Score: 0.8645833333333334

	Model	Score
0	LogisticRegression	0.872917
1	SVC	0.868750
2	DecisionTreeClassifier	0.864583

Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
model2 = RandomForestClassifier(random_state=1)
model2.fit(X_train, Y_train)
y_pred = model2.predict(X_test)

from sklearn.metrics import accuracy_score
print("Accuracy Score:", accuracy_score(Y_test, y_pred))
model_res.loc[len(model_res)] = ['RandomForestClassifier',
accuracy_score(Y_test, y_pred)]
model_res
```

Accuracy Score: 0.89375

	Model	Score
0	LogisticRegression	0.872917
1	SVC	0.868750
2	DecisionTreeClassifier	0.864583
3	RandomForestClassifier	0.893750

```
model_res = model_res.sort_values(by='Score', ascending=False)
model_res
```

	Model	Score
3	RandomForestClassifier	0.893750
0	LogisticRegression	0.872917
1	SVC	0.868750
2	DecisionTreeClassifier	0.864583